

Quora_Similarity_Case_Study

January 20, 2019

```
In [1]: import pandas as pd
import matplotlib.pyplot as plt
import re
import time
import warnings
import numpy as np
from nltk.corpus import stopwords
from sklearn.preprocessing import normalize
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.feature_extraction.text import TfidfVectorizer
warnings.filterwarnings("ignore")
import sys
import os
import pandas as pd
import numpy as np
from tqdm import tqdm
from sklearn.calibration import CalibratedClassifierCV

from sklearn.linear_model import SGDClassifier
from sklearn.metrics import confusion_matrix
from sklearn.metrics.classification import accuracy_score, log_loss
import seaborn as sns

In [2]: # avoid decoding problems
df = pd.read_csv("F:/Applied AI Course/quora similarity/train.csv")

# encode questions to unicode
# https://stackoverflow.com/a/6812069
# ----- python 2 -----
# df['question1'] = df['question1'].apply(lambda x: unicode(str(x), "utf-8"))
# df['question2'] = df['question2'].apply(lambda x: unicode(str(x), "utf-8"))
# ----- python 3 -----
df['question1'] = df['question1'].apply(lambda x: str(x))
df['question2'] = df['question2'].apply(lambda x: str(x))

In [3]: df.head()

Out[3]:
```

	id	qid1	qid2	question1	
0	0	1	2	What is the step by step guide to invest in sh...	\

```

1  1      3      4  What is the story of Kohinoor (Koh-i-Noor) Dia...
2  2      5      6  How can I increase the speed of my internet co...
3  3      7      8  Why am I mentally very lonely? How can I solve...
4  4      9     10  Which one dissolve in water quickly sugar, salt...

```

```

                                question2  is_duplicate
0  What is the step by step guide to invest in sh...      0
1  What would happen if the Indian government sto...      0
2  How can Internet speed be increased by hacking...      0
3  Find the remainder when  $23^{24}$  i...      0
4                Which fish would survive in salt water?      0

```

```

In [3]: from sklearn.feature_extraction.text import TfidfVectorizer
        from sklearn.feature_extraction.text import CountVectorizer

```

```
tfidf = TfidfVectorizer(max_features=500)
```

```

In [11]: # merge texts
         questions1 = (df['question1'].values)
         questions2 = (df['question2'].values)
         data_q1=(tfidf.fit_transform(questions1))
         data_q2=(tfidf.fit_transform(questions2))

```

```
In [5]: data_q1.get_shape()
```

```
Out[5]: (404290, 500)
```

```
In [12]: data_q1=data_q1.toarray()
```

```
In [13]: data_q2=data_q2.toarray()
```

```
In [14]: data_q1=pd.DataFrame(data_q1)
```

```
In [15]: data_q2=pd.DataFrame(data_q2)
```

```
In [16]: data_q1.head()
```

```

Out[16]:    0      1      2      3      4      5      6      7      8      9      ...    490  491  492  493  \
0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  ...  0.0  0.0  0.0  0.0
1  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  ...  0.0  0.0  0.0  0.0
2  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  ...  0.0  0.0  0.0  0.0
3  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  ...  0.0  0.0  0.0  0.0
4  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  0.0  ...  0.0  0.0  0.0  0.0

      494  495  496  497  498  499
0  0.0  0.0  0.0  0.0  0.0  0.0
1  0.0  0.0  0.0  0.0  0.0  0.0
2  0.0  0.0  0.0  0.0  0.0  0.0

```

```

3  0.0  0.0  0.0  0.0  0.0  0.0
4  0.0  0.0  0.0  0.0  0.0  0.0

```

[5 rows x 500 columns]

```

In [17]: #prepro_features_train.csv (Simple Preprocessing Feartures)
#nlp_features_train.csv (NLP Features)
if os.path.isfile('F:/Applied AI Course/quora similarity/nlp_features_train.csv'):
    dfnlp = pd.read_csv("F:/Applied AI Course/quora similarity/nlp_features_train.csv")
else:
    print("download nlp_features_train.csv from drive or run previous notebook")

if os.path.isfile('F:/Applied AI Course/quora similarity/df_fe_without_preprocessing_train.csv'):
    dfppro = pd.read_csv("F:/Applied AI Course/quora similarity/df_fe_without_preprocessing_train.csv")
else:
    print("download df_fe_without_preprocessing_train.csv from drive or run previous notebook")

In [18]: df1 = dfnlp.drop(['qid1', 'qid2', 'question1', 'question2'], axis=1)
df2 = dfppro.drop(['qid1', 'qid2', 'question1', 'question2', 'is_duplicate'], axis=1)
df3 = df.drop(['qid1', 'qid2', 'question1', 'question2', 'is_duplicate'], axis=1)

```

In [12]: df1.head()

```

Out[12]:   id  is_duplicate  cwc_min  cwc_max  csc_min  csc_max  ctc_min  \
0    0             0  0.999980  0.833319  0.999983  0.999983  0.916659
1    1             0  0.799984  0.399996  0.749981  0.599988  0.699993
2    2             0  0.399992  0.333328  0.399992  0.249997  0.399996
3    3             0  0.000000  0.000000  0.000000  0.000000  0.000000
4    4             0  0.399992  0.199998  0.999950  0.666644  0.571420

      ctc_max  last_word_eq  first_word_eq  abs_len_diff  mean_len  \
0  0.785709             0.0             1.0             2.0      13.0
1  0.466664             0.0             1.0             5.0      12.5
2  0.285712             0.0             1.0             4.0      12.0
3  0.000000             0.0             0.0             2.0      12.0
4  0.307690             0.0             1.0             6.0      10.0

      token_set_ratio  token_sort_ratio  fuzz_ratio  fuzz_partial_ratio  \
0                  100                 93         93                 100
1                   86                 63         66                 75
2                   66                 66         54                 54
3                   36                 36         35                 40
4                   67                 47         46                 56

      longest_substr_ratio
0          0.982759
1          0.596154
2          0.166667

```

```

3          0.039216
4          0.175000

```

```
In [13]: df2.head()
```

```

Out[13]:   id  freq_qid1  freq_qid2  q1len  q2len  q1_n_words  q2_n_words  \
0    0         1         1      66      57         14         12
1    1         4         1      51      88          8         13
2    2         1         1      73      59         14         10
3    3         1         1      50      65         11          9
4    4         3         1      76      39         13          7

      word_Common  word_Total  word_share  freq_q1+q2  freq_q1-q2
0             10.0         23.0    0.434783          2          0
1              4.0         20.0    0.200000          5          3
2              4.0         24.0    0.166667          2          0
3              0.0         19.0    0.000000          2          0
4              2.0         20.0    0.100000          4          2

```

```

In [14]: print("Number of features in nlp dataframe :", df1.shape[1])
print("Number of features in preprocessed dataframe :", df2.shape[1])
print("Number of features in question1 tf-idf dataframe :", data_q1.shape[1])
print("Number of features in question2 tf-idf dataframe :", data_q2.shape[1])
print("Number of features in final dataframe :", df1.shape[1]+df2.shape[1]+data_q1.shape[1]+data_q2.shape[1])

```

```

Number of features in nlp dataframe : 17
Number of features in preprocessed dataframe : 12
Number of features in question1 tf-idf dataframe : 500
Number of features in question2 tf-idf dataframe : 500
Number of features in final dataframe : 1029

```

```

In [21]: # storing the final features to csv file
from sklearn.utils import resample
df1=resample(df1, n_samples=100000, random_state=30)
df2=resample(df2, n_samples=100000, random_state=30)
data_q1=resample(data_q1, n_samples=100000, random_state=30)
data_q2=resample(data_q2, n_samples=100000, random_state=30)

```

```
In [22]: print("Number of datapoints in final dataframe :", df1.shape[0]+df2.shape[0]+data_q1.shape[0]+data_q2.shape[0])
```

```
Number of datapoints in final dataframe : 400000
```

```

In [27]: data_q1['id']=df1['id']
data_q2['id']=df1['id']

```

```

In [24]: df1 = df1.merge(df2, on='id',how='left')
df2 = data_q1.merge(data_q2, on='id',how='left')
result = df1.merge(df2, on='id',how='left')

```

```
In [25]: result.shape
```

```
Out[25]: (307762, 1039)
```

```
In [26]: result=resample(result, n_samples=100000, random_state=30)
```

```
In [28]: result.to_csv('final_features.csv')
```

```
In [4]: result=pd.read_csv('final_features.csv')
```

```
In [5]: result.shape
```

```
Out[5]: (100000, 1040)
```

```
In [7]: y_true = result['is_duplicate']
```

```
In [9]: result.drop([ 'id','is_duplicate'], axis=1, inplace=True)
```

Random train test split(70:30)

```
In [10]: from sklearn.model_selection import train_test_split
         X_train,X_test, y_train, y_test = train_test_split(result, y_true, stratify=y_true, t
```

```
In [11]: print("Number of data points in train data :",X_train.shape)
         print("Number of data points in test data :",X_test.shape)
```

Number of data points in train data : (70000, 1038)

Number of data points in test data : (30000, 1038)

```
In [12]: # This function plots the confusion matrices given y_i, y_i_hat.
def plot_confusion_matrix(test_y, predict_y):
    C = confusion_matrix(test_y, predict_y)
    # C = 9,9 matrix, each cell (i,j) represents number of points of class i are pred

    A = (((C.T)/(C.sum(axis=1))).T)
    #divid each element of the confusion matrix with the sum of elements in that colu

    # C = [[1, 2],
    #       [3, 4]]
    # C.T = [[1, 3],
    #         [2, 4]]
    # C.sum(axis = 1) axis=0 corresponds to columns and axis=1 corresponds to rows in
    # C.sum(axix =1) = [[3, 7]]
    # ((C.T)/(C.sum(axis=1))) = [[1/3, 3/7]
    #                          [2/3, 4/7]]

    # ((C.T)/(C.sum(axis=1))).T = [[1/3, 2/3]
    #                          [3/7, 4/7]]
    # sum of row elements = 1
```

```

B=(C/C.sum(axis=0))
#divid each element of the confusion matrix with the sum of elements in that row
# C = [[1, 2],
#      [3, 4]]
# C.sum(axis = 0) axis=0 corresonds to columns and axis=1 corresponds to rows in
# C.sum(axix =0) = [[4, 6]]
# (C/C.sum(axis=0)) = [[1/4, 2/6],
#                      [3/4, 4/6]]
plt.figure(figsize=(20,4))

labels = [1,2]
# representing A in heatmap format
cmap=sns.light_palette("blue")
plt.subplot(1, 3, 1)
sns.heatmap(C, annot=True, cmap=cmap, fmt=".3f", xticklabels=labels, yticklabels=labels)
plt.xlabel('Predicted Class')
plt.ylabel('Original Class')
plt.title("Confusion matrix")

plt.subplot(1, 3, 2)
sns.heatmap(B, annot=True, cmap=cmap, fmt=".3f", xticklabels=labels, yticklabels=labels)
plt.xlabel('Predicted Class')
plt.ylabel('Original Class')
plt.title("Precision matrix")

plt.subplot(1, 3, 3)
# representing B in heatmap format
sns.heatmap(A, annot=True, cmap=cmap, fmt=".3f", xticklabels=labels, yticklabels=labels)
plt.xlabel('Predicted Class')
plt.ylabel('Original Class')
plt.title("Recall matrix")

plt.show()

```

Logistic Regression with hyperparameter tuning

```

In [48]: alpha = [10 ** x for x in range(-5, 2)] # hyperparam for SGD classifier.

# read more about SGDClassifier() at http://scikit-learn.org/stable/modules/generated/sklearn.linear_model.SGDClassifier.html
# -----
# default parameters
# SGDClassifier(loss=hinge, penalty=l2, alpha=0.0001, l1_ratio=0.15, fit_intercept=True,
# shuffle=True, verbose=0, epsilon=0.1, n_jobs=1, random_state=None, learning_rate=optimal,
# class_weight=None, warm_start=False, average=False, n_iter=None)

# some of methods
# fit(X, y[, coef_init, intercept_init, ])          Fit linear model with Stochastic Gradient Descent

```

```

# predict(X)          Predict class labels for samples in X.

#-----
# video link:
#-----

log_error_array=[]
for i in alpha:
    clf = SGDClassifier(alpha=i, penalty='l2', loss='log', random_state=42)
    clf.fit(X_train, y_train)
    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(X_train, y_train)
    predict_y = sig_clf.predict_proba(X_test)
    log_error_array.append(log_loss(y_test, predict_y, labels=clf.classes_, eps=1e-15))
    print('For values of alpha = ', i, "The log loss is:", log_loss(y_test, predict_y,

fig, ax = plt.subplots()
ax.plot(alpha, log_error_array, c='g')
for i, txt in enumerate(np.round(log_error_array, 3)):
    ax.annotate((alpha[i], np.round(txt, 3)), (alpha[i], log_error_array[i]))
plt.grid()
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.show()

best_alpha = np.argmin(log_error_array)
clf = SGDClassifier(alpha=alpha[best_alpha], penalty='l2', loss='log', random_state=42)
clf.fit(X_train, y_train)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(X_train, y_train)

predict_y = sig_clf.predict_proba(X_train)
print('For values of best alpha = ', alpha[best_alpha], "The train log loss is:", log_loss(y_train, predict_y, labels=clf.classes_, eps=1e-15))
predict_y = sig_clf.predict_proba(X_test)
print('For values of best alpha = ', alpha[best_alpha], "The test log loss is:", log_loss(y_test, predict_y, labels=clf.classes_, eps=1e-15))
predicted_y = np.argmax(predict_y, axis=1)
print("Total number of data points :", len(predicted_y))
plot_confusion_matrix(y_test, predicted_y)

```

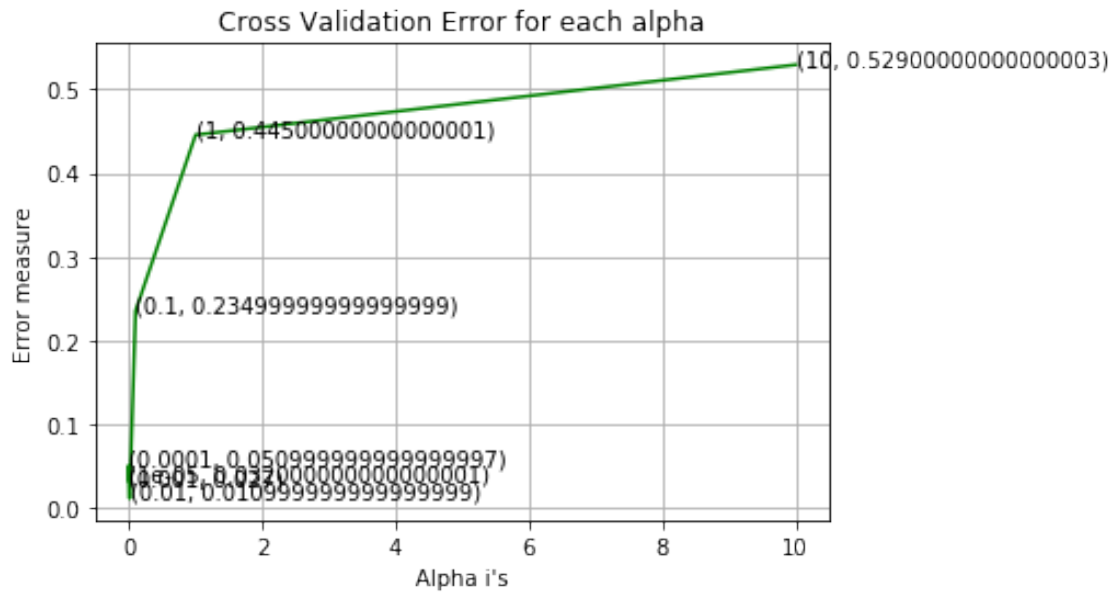
```

For values of alpha = 1e-05 The log loss is: 0.0324060901258
For values of alpha = 0.0001 The log loss is: 0.051059843091
For values of alpha = 0.001 The log loss is: 0.0265776544546
For values of alpha = 0.01 The log loss is: 0.011433874772
For values of alpha = 0.1 The log loss is: 0.235486697939
For values of alpha = 1 The log loss is: 0.445421194718

```

For values of alpha = 10 The log loss is: 0.52910987166

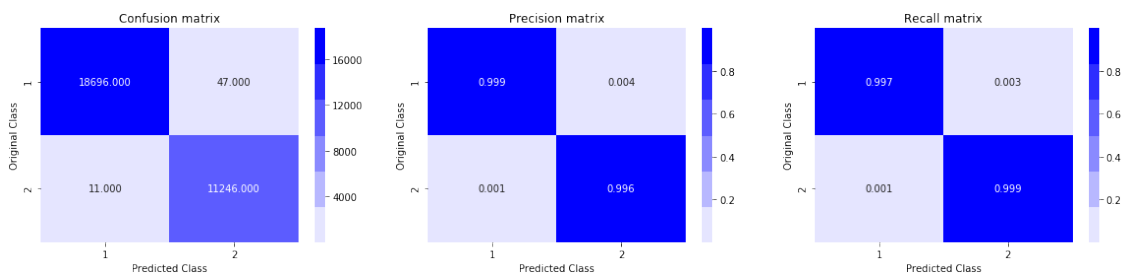
<Figure size 1440x288 with 0 Axes>



For values of best alpha = 0.01 The train log loss is: 0.0108237418247

For values of best alpha = 0.01 The test log loss is: 0.011433874772

Total number of data points : 30000



Linear SVM with hyperparameter tuning

```
In [49]: alpha = [10 ** x for x in range(-5, 2)] # hyperparam for SGD classifier.
```

```
# read more about SGDClassifier() at http://scikit-learn.org/stable/modules/generated,
# -----
```



```

# default parameters
# SGDClassifier(loss=hinge, penalty=l2, alpha=0.0001, l1_ratio=0.15, fit_intercept=True,
# shuffle=True, verbose=0, epsilon=0.1, n_jobs=1, random_state=None, learning_rate=optimal,
# class_weight=None, warm_start=False, average=False, n_iter=None)

# some of methods
# fit(X, y[, coef_init, intercept_init, ])          Fit linear model with Stochastic Gradient Descent
# predict(X)          Predict class labels for samples in X.

#-----
# video link:
#-----

log_error_array=[]
for i in alpha:
    clf = SGDClassifier(alpha=i, penalty='l1', loss='hinge', random_state=42)
    clf.fit(X_train, y_train)
    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(X_train, y_train)
    predict_y = sig_clf.predict_proba(X_test)
    log_error_array.append(log_loss(y_test, predict_y, labels=clf.classes_, eps=1e-15))
    print('For values of alpha = ', i, "The log loss is:", log_loss(y_test, predict_y, labels=clf.classes_, eps=1e-15))

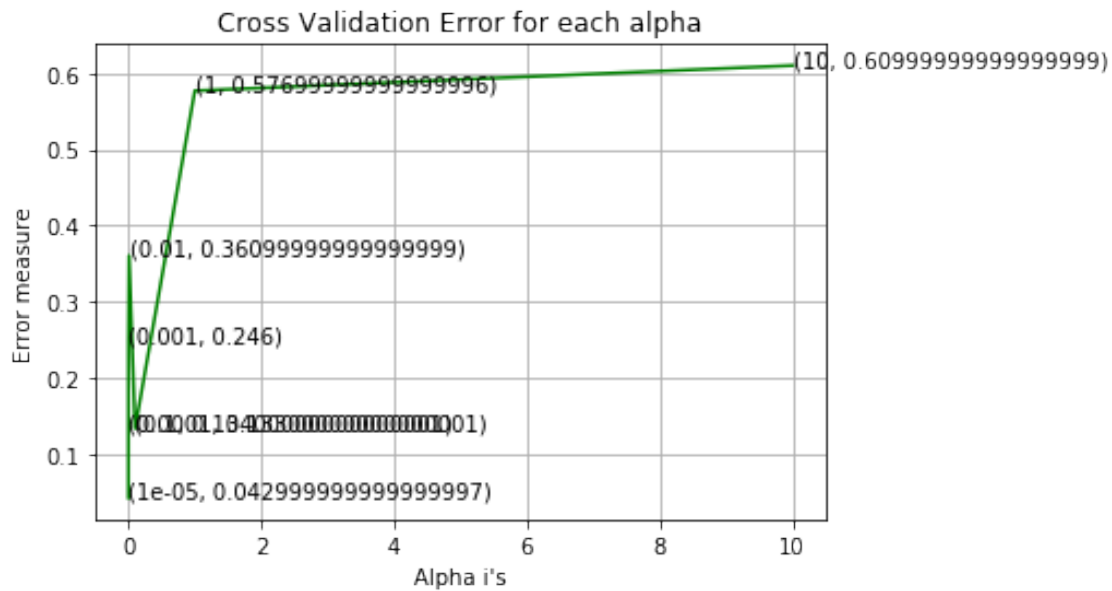
fig, ax = plt.subplots()
ax.plot(alpha, log_error_array, c='g')
for i, txt in enumerate(np.round(log_error_array, 3)):
    ax.annotate((alpha[i], np.round(txt, 3)), (alpha[i], log_error_array[i]))
plt.grid()
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.show()

best_alpha = np.argmin(log_error_array)
clf = SGDClassifier(alpha=alpha[best_alpha], penalty='l1', loss='hinge', random_state=42)
clf.fit(X_train, y_train)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(X_train, y_train)

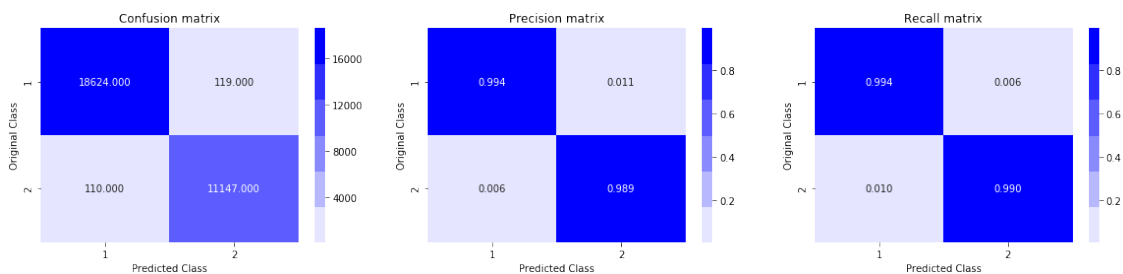
predict_y = sig_clf.predict_proba(X_train)
print('For values of best alpha = ', alpha[best_alpha], "The train log loss is:", log_loss(y_train, predict_y, labels=clf.classes_, eps=1e-15))
predict_y = sig_clf.predict_proba(X_test)
print('For values of best alpha = ', alpha[best_alpha], "The test log loss is:", log_loss(y_test, predict_y, labels=clf.classes_, eps=1e-15))
predicted_y = np.argmax(predict_y, axis=1)
print("Total number of data points :", len(predicted_y))
plot_confusion_matrix(y_test, predicted_y)

```

For values of alpha = 1e-05 The log loss is: 0.0427410657654
 For values of alpha = 0.0001 The log loss is: 0.133065455907
 For values of alpha = 0.001 The log loss is: 0.246352507555
 For values of alpha = 0.01 The log loss is: 0.36060485994
 For values of alpha = 0.1 The log loss is: 0.133812754686
 For values of alpha = 1 The log loss is: 0.576707543599
 For values of alpha = 10 The log loss is: 0.609841792942



For values of best alpha = 1e-05 The train log loss is: 0.0408144563488
 For values of best alpha = 1e-05 The test log loss is: 0.0427410657654
 Total number of data points : 30000



XGBoost

```
In [13]: from sklearn.model_selection import RandomizedSearchCV
         from scipy import stats
```

```

from xgboost import XGBClassifier
import xgboost as xgb

In [14]: param_dist = {'n_estimators': [25,50,75,100],
                        'learning_rate': stats.uniform(0.01, 0.07),
                        'max_depth': [5,7,9,11]
                      }

model=RandomizedSearchCV(XGBClassifier(n_jobs=-1), param_distributions=param_dist, sc
model.fit(X_train, y_train)
print(model.best_estimator_)
print(model.score(X_test, y_test))

XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=1,
               colsample_bytree=1, gamma=0, learning_rate=0.068660486889955685,
               max_delta_step=0, max_depth=9, min_child_weight=1, missing=None,
               n_estimators=100, n_jobs=-1, nthread=None,
               objective='binary:logistic', random_state=0, reg_alpha=0,
               reg_lambda=1, scale_pos_weight=1, seed=None, silent=True,
               subsample=1)
-0.227220311349

In [15]: params={}
         params['n_estimators'] = 100
         params['learning_rate'] = 0.068
         params['max_depth'] = 9

In [16]: d_train = xgb.DMatrix(X_train, label=y_train)
         d_test = xgb.DMatrix(X_test, label=y_test)

         watchlist = [(d_train, 'train'), (d_test, 'valid')]

         bst = xgb.train(params, d_train, 400, watchlist, early_stopping_rounds=20, verbose_ev

         xgdmatrix = xgb.DMatrix(X_train,y_train)
         predict_y = bst.predict(d_train)
         print("The train log loss is:",log_loss(y_train, predict_y, labels=clf.classes_, eps=1e-
         predict_y = bst.predict(d_test)
         print("The test log loss is:",log_loss(y_test, predict_y, labels=clf.classes_, eps=1e-

[18:50:37] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 492 extra nodes, 0 pruned n
[0]      train-rmse:0.479837      valid-rmse:0.480244
Multiple eval metrics have been passed: 'valid-rmse' will be used for early stopping.

Will train until valid-rmse hasn't improved in 20 rounds.
[18:50:40] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 452 extra nodes, 0 pruned n
[18:50:42] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 518 extra nodes, 0 pruned n
[18:50:44] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 504 extra nodes, 0 pruned n
[18:50:46] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 484 extra nodes, 0 pruned n

```



```

[18:52:26] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 340 extra nodes, 0 pruned n
[18:52:29] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 362 extra nodes, 0 pruned n
[50]      train-rmse:0.244683      valid-rmse:0.263502
[18:52:31] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 338 extra nodes, 0 pruned n
[18:52:33] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 208 extra nodes, 0 pruned n
[18:52:35] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 246 extra nodes, 0 pruned n
[18:52:37] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 210 extra nodes, 0 pruned n
[18:52:39] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 258 extra nodes, 0 pruned n
[18:52:41] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 222 extra nodes, 0 pruned n
[18:52:43] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 170 extra nodes, 0 pruned n
[18:52:45] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 214 extra nodes, 0 pruned n
[18:52:47] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 216 extra nodes, 0 pruned n
[18:52:49] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 248 extra nodes, 0 pruned n
[60]      train-rmse:0.239415      valid-rmse:0.259318
[18:52:51] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 180 extra nodes, 0 pruned n
[18:52:53] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 218 extra nodes, 0 pruned n
[18:52:55] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 184 extra nodes, 0 pruned n
[18:52:57] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 190 extra nodes, 0 pruned n
[18:52:59] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 248 extra nodes, 0 pruned n
[18:53:01] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 164 extra nodes, 0 pruned n
[18:53:04] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 254 extra nodes, 0 pruned n
[18:53:06] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 162 extra nodes, 0 pruned n
[18:53:08] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 266 extra nodes, 0 pruned n
[18:53:10] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 164 extra nodes, 0 pruned n
[70]      train-rmse:0.23539      valid-rmse:0.256351
[18:53:12] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 168 extra nodes, 0 pruned n
[18:53:14] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 196 extra nodes, 0 pruned n
[18:53:16] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 124 extra nodes, 0 pruned n
[18:53:19] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 166 extra nodes, 0 pruned n
[18:53:21] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 140 extra nodes, 0 pruned n
[18:53:23] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 202 extra nodes, 0 pruned n
[18:53:25] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 188 extra nodes, 0 pruned n
[18:53:28] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 224 extra nodes, 0 pruned n
[18:53:30] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 140 extra nodes, 0 pruned n
[18:53:33] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 198 extra nodes, 0 pruned n
[80]      train-rmse:0.232745      valid-rmse:0.254318
[18:53:35] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 90 extra nodes, 0 pruned n
[18:53:37] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 144 extra nodes, 0 pruned n
[18:53:39] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 232 extra nodes, 0 pruned n
[18:53:41] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 112 extra nodes, 0 pruned n
[18:53:43] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 176 extra nodes, 0 pruned n
[18:53:46] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 108 extra nodes, 0 pruned n
[18:53:48] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 114 extra nodes, 0 pruned n
[18:53:50] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 218 extra nodes, 0 pruned n
[18:53:52] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 196 extra nodes, 0 pruned n
[18:53:54] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 150 extra nodes, 0 pruned n
[90]      train-rmse:0.230712      valid-rmse:0.252827
[18:53:57] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 184 extra nodes, 0 pruned n

```


[19:01:48]	src/tree/updater_prune.cc:74:	tree pruning end,	1 roots,	206 extra nodes,	0 pruned n
[19:01:50]	src/tree/updater_prune.cc:74:	tree pruning end,	1 roots,	72 extra nodes,	0 pruned n
[19:01:52]	src/tree/updater_prune.cc:74:	tree pruning end,	1 roots,	66 extra nodes,	0 pruned n
[19:01:54]	src/tree/updater_prune.cc:74:	tree pruning end,	1 roots,	162 extra nodes,	0 pruned n
[19:01:56]	src/tree/updater_prune.cc:74:	tree pruning end,	1 roots,	164 extra nodes,	0 pruned n
[19:01:58]	src/tree/updater_prune.cc:74:	tree pruning end,	1 roots,	122 extra nodes,	0 pruned n
[19:02:01]	src/tree/updater_prune.cc:74:	tree pruning end,	1 roots,	160 extra nodes,	0 pruned n
[19:02:03]	src/tree/updater_prune.cc:74:	tree pruning end,	1 roots,	146 extra nodes,	0 pruned n
[19:02:05]	src/tree/updater_prune.cc:74:	tree pruning end,	1 roots,	76 extra nodes,	0 pruned n
[19:02:07]	src/tree/updater_prune.cc:74:	tree pruning end,	1 roots,	124 extra nodes,	0 pruned n
[320]	train-rmse:0.196851	valid-rmse:0.228969			
[19:02:09]	src/tree/updater_prune.cc:74:	tree pruning end,	1 roots,	166 extra nodes,	0 pruned n
[19:02:11]	src/tree/updater_prune.cc:74:	tree pruning end,	1 roots,	92 extra nodes,	0 pruned n
[19:02:13]	src/tree/updater_prune.cc:74:	tree pruning end,	1 roots,	90 extra nodes,	0 pruned n
[19:02:15]	src/tree/updater_prune.cc:74:	tree pruning end,	1 roots,	96 extra nodes,	0 pruned n
[19:02:17]	src/tree/updater_prune.cc:74:	tree pruning end,	1 roots,	104 extra nodes,	0 pruned n
[19:02:19]	src/tree/updater_prune.cc:74:	tree pruning end,	1 roots,	104 extra nodes,	0 pruned n
[19:02:21]	src/tree/updater_prune.cc:74:	tree pruning end,	1 roots,	32 extra nodes,	0 pruned n
[19:02:24]	src/tree/updater_prune.cc:74:	tree pruning end,	1 roots,	142 extra nodes,	0 pruned n
[19:02:26]	src/tree/updater_prune.cc:74:	tree pruning end,	1 roots,	190 extra nodes,	0 pruned n
[19:02:28]	src/tree/updater_prune.cc:74:	tree pruning end,	1 roots,	114 extra nodes,	0 pruned n
[330]	train-rmse:0.194551	valid-rmse:0.227301			
[19:02:30]	src/tree/updater_prune.cc:74:	tree pruning end,	1 roots,	96 extra nodes,	0 pruned n
[19:02:32]	src/tree/updater_prune.cc:74:	tree pruning end,	1 roots,	96 extra nodes,	0 pruned n
[19:02:34]	src/tree/updater_prune.cc:74:	tree pruning end,	1 roots,	144 extra nodes,	0 pruned n
[19:02:36]	src/tree/updater_prune.cc:74:	tree pruning end,	1 roots,	128 extra nodes,	0 pruned n
[19:02:38]	src/tree/updater_prune.cc:74:	tree pruning end,	1 roots,	104 extra nodes,	0 pruned n
[19:02:40]	src/tree/updater_prune.cc:74:	tree pruning end,	1 roots,	166 extra nodes,	0 pruned n
[19:02:42]	src/tree/updater_prune.cc:74:	tree pruning end,	1 roots,	84 extra nodes,	0 pruned n
[19:02:44]	src/tree/updater_prune.cc:74:	tree pruning end,	1 roots,	80 extra nodes,	0 pruned n
[19:02:46]	src/tree/updater_prune.cc:74:	tree pruning end,	1 roots,	110 extra nodes,	0 pruned n
[19:02:48]	src/tree/updater_prune.cc:74:	tree pruning end,	1 roots,	66 extra nodes,	0 pruned n
[340]	train-rmse:0.193554	valid-rmse:0.226623			
[19:02:50]	src/tree/updater_prune.cc:74:	tree pruning end,	1 roots,	138 extra nodes,	0 pruned n
[19:02:52]	src/tree/updater_prune.cc:74:	tree pruning end,	1 roots,	142 extra nodes,	0 pruned n
[19:02:54]	src/tree/updater_prune.cc:74:	tree pruning end,	1 roots,	212 extra nodes,	0 pruned n
[19:02:57]	src/tree/updater_prune.cc:74:	tree pruning end,	1 roots,	130 extra nodes,	0 pruned n
[19:02:59]	src/tree/updater_prune.cc:74:	tree pruning end,	1 roots,	90 extra nodes,	0 pruned n
[19:03:01]	src/tree/updater_prune.cc:74:	tree pruning end,	1 roots,	120 extra nodes,	0 pruned n
[19:03:03]	src/tree/updater_prune.cc:74:	tree pruning end,	1 roots,	118 extra nodes,	0 pruned n
[19:03:05]	src/tree/updater_prune.cc:74:	tree pruning end,	1 roots,	158 extra nodes,	0 pruned n
[19:03:07]	src/tree/updater_prune.cc:74:	tree pruning end,	1 roots,	58 extra nodes,	0 pruned n
[19:03:09]	src/tree/updater_prune.cc:74:	tree pruning end,	1 roots,	158 extra nodes,	0 pruned n
[350]	train-rmse:0.192229	valid-rmse:0.225766			
[19:03:11]	src/tree/updater_prune.cc:74:	tree pruning end,	1 roots,	140 extra nodes,	0 pruned n
[19:03:13]	src/tree/updater_prune.cc:74:	tree pruning end,	1 roots,	110 extra nodes,	0 pruned n
[19:03:15]	src/tree/updater_prune.cc:74:	tree pruning end,	1 roots,	150 extra nodes,	0 pruned n


```
[19:04:47] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 152 extra nodes, 0 pruned nodes
[19:04:49] src/tree/updater_prune.cc:74: tree pruning end, 1 roots, 52 extra nodes, 0 pruned nodes
[399]          train-rmse:0.187477          valid-rmse:0.22261
```

```
-----

NameError                                Traceback (most recent call last)
```

```
<ipython-input-16-d4a006c3671e> in <module>()
      8 xgdmatrix = xgb.DMatrix(X_train,y_train)
      9 predict_y = bst.predict(d_train)
----> 10 print("The train log loss is:",log_loss(y_train, predict_y, labels=clf.classes_, epsilon=0.0001))
      11 predict_y = bst.predict(d_test)
      12 print("The test log loss is:",log_loss(y_test, predict_y, labels=clf.classes_, epsilon=0.0001))
```

```
NameError: name 'log_loss' is not defined
```

```
In [17]: print("The train log loss is:",0.187)
```

```
print("The test log loss is:",0.222)
```

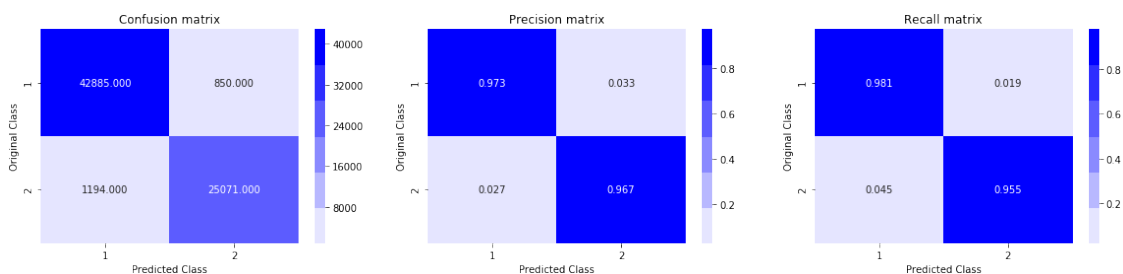
```
The train log loss is: 0.187
```

```
The test log loss is: 0.222
```

```
In [22]: predicted_y = np.array(predict_y>0.5,dtype=int)
print("Total number of data points :", len(predicted_y))
plot_confusion_matrix(y_train, predicted_y)
```

```
Total number of data points : 70000
```

```
<Figure size 1440x288 with 0 Axes>
```



Conclusion

```
In [23]: from prettytable import PrettyTable
x = PrettyTable()

x.field_names = ["Model ", "Train Loss", "Test Loss"]

x.add_row(['Logistic Regression',0.010,0.011])
x.add_row(['Linear SVM',0.040,0.042])
x.add_row(['XGBoost',0.187,0.222])

print(x)
```

Model	Train Loss	Test Loss
Logistic Regression	0.01	0.011
Linear SVM	0.04	0.042
XGBoost	0.187	0.222